

DASC 41103
Machine Learning
Project 1
Foundations of Classification Algorithms

Total Points: 200

Lesson Objective Alignment:

- Design and implement machine learning models using Python.
 - Develop machine learning algorithms for practical applications
 - Understand and implement perceptron and Adaline algorithms.
 - Apply logistic regression and support vector machines using scikit-learn.
- Evaluate model performance and decision boundaries.
- Work with real-world datasets and apply feature preprocessing.
- Communicate machine learning principles and methods to diverse audiences.

Project Objective: Demonstrate ability to apply simple machine learning classification algorithms in Python.

Description: For this project, students will work in **pairs (groups of 2)** to complete each part, provide requested deliverables, and answer any questions.

Data: Use the provided data, which comes from the **UCI Adult Income dataset** (also known as the Census Income dataset) to predict whether a person earns more than \$50K/year.

- project_adult.csv
- project_validation_inputs.csv

You can learn more about the data variables by going to <https://archive.ics.uci.edu/dataset/2/adult>.

Parts:

1. Preprocess the dataset: **project_adult.csv** & **project_validation_inputs.csv**
 - a. Handle missing values.
 - b. Encode categorical features.
 - c. Standardize numerical features.
2. Implement the **Perceptron** and **Adaline** algorithms
 - a. Train Perceptron and Adaline models (at least AdalineSGD).
 - b. Plot the number of misclassifications (Perceptron) and MSE (Adaline) over epochs.
 - c. Find the accuracy of your best models from both algorithms
 - d. Use best performing models to predict outputs for **project_validation_inputs**.

- e. Use provided code to implement scikit-learn's Perceptron and Adaline algorithms. Find the accuracy on the validation data using your best models for both.
3. Implement Logistic Regression and SVM using scikit-learn
 - a. Train Logistic Regression and SVM models using scikit-learn.
 - b. Find the accuracy of your best models from both algorithms
 - c. Use best performing models to predict outputs for **project_validation_inputs**.
 - d. Select 2 features and visualize the decision boundaries.
4. Reflection and Conceptual Questions
 - a. Why is feature scaling important for gradient-based algorithms?
 - b. Explain the difference between batch gradient descent and stochastic gradient descent.
 - c. Why does scikit-learn Perceptron and Adline algorithms outperform book code?
 - i. **Research and develop an informed answer.**
 - d. Compare the decision boundaries of logistic regression and SVM.
 - e. What is the role of regularization in preventing overfitting?
 - f. Vary the C values of the scikit-learn LogisticRegression and linear SVC models with [0.01, 1.0, 100.0]. **Discuss the impact.**

Deliverables:

Students should submit a video recording of their presentation, their slides, the predicted values of their model on the validation inputs, and a link to their open GitHub repository via blackboard.

- 15-minute presentation that at a high-level discussing what you did in parts 1 – 3 and the answers to questions in part 4 in more depth.
- 4 validation output files. This should be the predictive values your models made on project_validation_inputs. We will use these files to determine the accuracy based on the actual outputs. **This should be the scikit-learn implementations of each.** You should use the following names:
 - Group_#_Perceptron_PredictedOutputs.csv
 - Group_#_Adaline_PredictedOutputs.csv
 - Group_#_LogisticRegression_PredictedOutputs.csv
 - Group_#_SVM_PredictedOutputs.csv
- Link to your group's open GitHub repository with code that is clean and well-commented.

Grading:

All team members will receive the same grade unless a team member requests otherwise. Grades will be based on the grading rubric below.

Project Grading Rubric

| Group Number: | Points |
|--|--------|
| 1. Presentation: | |
| Project Discussion: 30 points: A clear, concise, and high-level summary of the methodology, challenges, and results for Parts 1, 2, and 3. Demonstrates a strong understanding of the project's practical steps. 15 points: Provides a partial summary of the project parts. Some key steps or findings are missing or unclear. 0 points: The project overview is missing or lacks a coherent narrative. | 30 |
| Conceptual Depth: 30 points: Offers a comprehensive and articulate explanation of the concepts from Part 4, demonstrating a deep understanding of algorithms and theory. 15 points: Provides a basic but incomplete explanation of the concepts. Lacks depth or contains some inaccuracies. 0 points: The conceptual answers are missing or incorrect. | 30 |
| Presentation Quality: 20 points: The presentation is well-structured, professional, and within the 15-minute time limit. Visuals on the slides are clean and easy to follow. Audio and video quality are excellent. 10 points: The presentation is somewhat disorganized or exceeds the time limit. Some visuals are cluttered, or the audio/video quality is poor. 0 points: The presentation is incomprehensible, significantly exceeds the time limit, or was not submitted. | 20 |
| 2. Code and Repository | |
| Code Functionality: 30 points: All code runs without errors, producing the expected outputs and plots. The implementations of Perceptron and Adaline are correct. 15 points: The code runs with minor errors or does not fully complete all tasks. 0 points: The code fails to run or contains significant errors, making it unusable. | 30 |
| Code Quality and Comments: 30 points: The code is clean, logically organized, and highly readable. It includes comprehensive and helpful comments explaining key functions, logical blocks, and complex steps. Variable names are descriptive. 15 points: The code is functional but could be cleaner. Comments are sparse or do not adequately explain the logic. 0 points: The code is unreadable, not commented, or disorganized. | 30 |
| GitHub Repository: 20 points: The repository is public, contains all required code and a README file, and is easily accessible via the link provided. The commit history is logical. 10 points: The repository is missing a README, is not public, or is difficult to navigate. 0 points: The GitHub link is broken, or the repository is not submitted. | 20 |

| | |
|---|----|
| 3. Model Performance and Deliverable | |
| Validation Output File Naming & Submission: 10 points: All four .csv files are submitted to Blackboard with the correct and exact file names as specified in the assignment. 5 points: Files are submitted but some have incorrect names or are missing. 0 points: No files are submitted. | 10 |
| Model Accuracy on Validation Set: 30 points: Your models achieve an acceptable level of accuracy on the provided validation set, demonstrating a robust implementation and an effective approach to feature preprocessing and model training. 15 points: Your models achieve a moderate level of accuracy. 0 points: Your models achieve a low level of accuracy, indicating fundamental errors in the implementation or training process. | 30 |
| Total | |